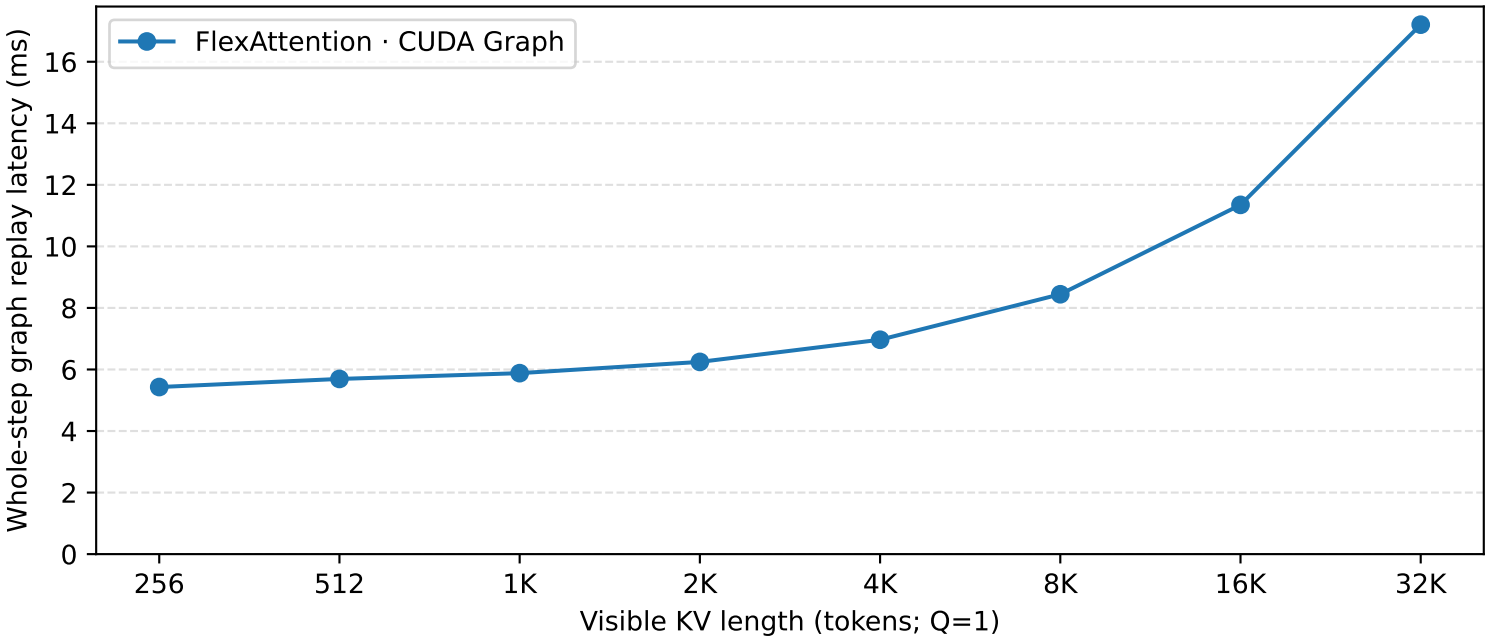
Decode · Llama-7B backbone · FP16 · batch 1
FlexAttention · CUDA Graph
Whole-step graph replay latency (ms)
0
2
4
6
8
10
12
14
16
256
512
1K
2K
4K
8K
16K
32K
Visible KV length (tokens; Q=1)