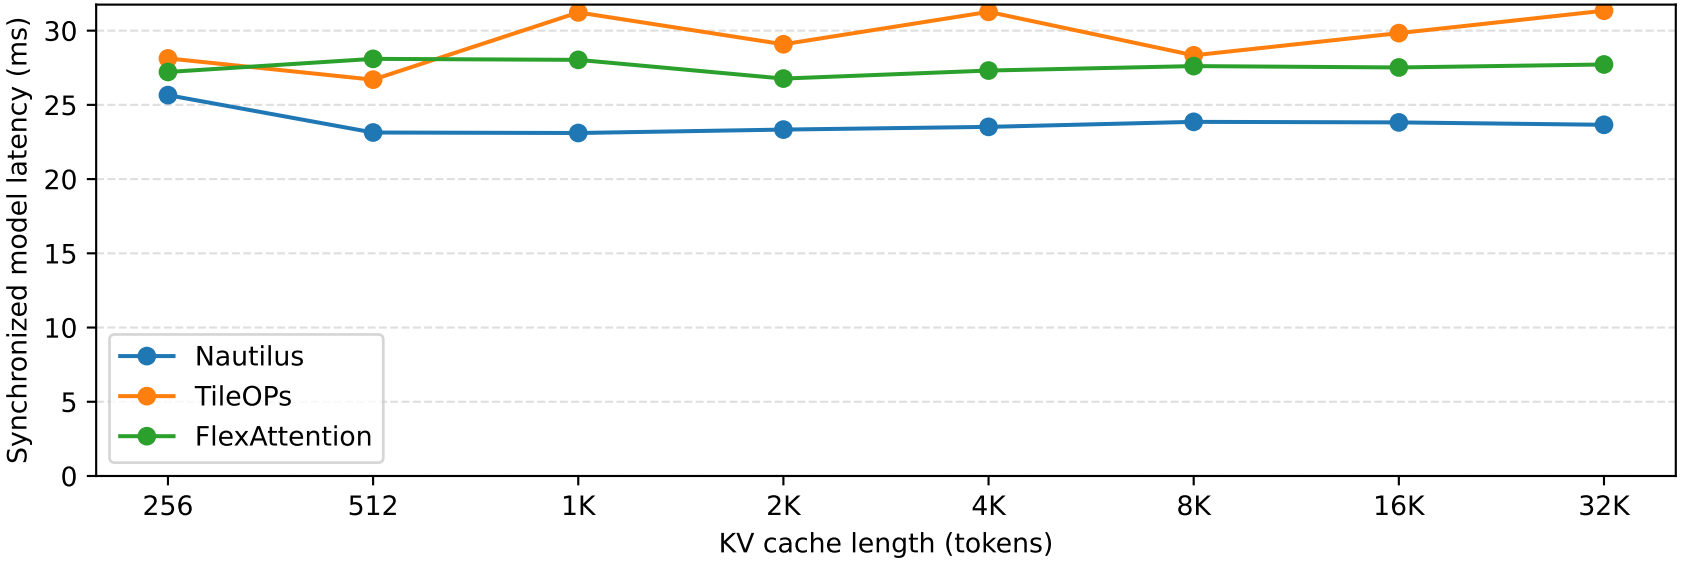
Decode · Llama-7B backbone · FP16 · batch 1 · query 1
Synchronized model latency (ms)
0
5
10
15
20
25
30
Nautilus
TileOPs
FlexAttention
256
512
1K
2K
4K
8K
16K
32K
KV cache length (tokens)